Ove Bepari

Site Reliability & Performance Engineer · Dhaka, Bangladesh

<redacted> / <redacted> / <redacted> / <redacted>

Site reliability engineer with **4 years** of experience operating a US electronic health records platform of 30+ JVM services handling millions of requests a day. With a background in backend development, I trace production issues through the JVM, the Linux kernel, and the network stack to their root cause, and implement lasting fixes.

01 Experience

Sep 2022 – Present

Application Operations & Reliability Engineer II

<redacted>

Part of <redacted> — electronic health records for the US disability services sector

- **Release Engineering:** Run production releases (canaries, rollbacks) for **30+** JVM and Docker services handling **millions of daily requests**. Automated them with Jenkins and Bash/Python scripts that deploy in dependency order (a DAG sorted with tsort) and smoke-test every rollout.
- **ML Infrastructure:** Deploy the team's vLLM inference service (Docker, NVIDIA GPU runtime) on company-owned H200 servers, serving an off-the-shelf Gemma 4 26B-A4B model at **~100k requests/day** for summarization, analysis, and documentation features. Maintain the JVM client service that connects the SaaS app to it, instrumented with Datadog APM.
- **Observability:** Rolled out Datadog APM across **40+** services and built its dashboards and monitors. Co-architected the log pipeline (Vector → Kafka → Falcon LogScale) and set up its Kafka and LogScale clusters. Cut MTTD from days to **under 5 minutes**.
- **Production Debugging & On-call:** Take shifts on the team's SRE on-call. Debug production issues with JFR, Wireshark, and Linux tooling until the real cause turns up, from file-descriptor leaks to exhausted thread pools and page-cache thrashing. Most are fixed within **1–2 hours**.
- **Distributed Systems:** Traced split-brain, outages, and data inconsistency in our Hazelcast clusters to a faulty leader-election mechanism, and proposed fixes that the developers implemented. Also look after Oracle Coherence clusters in production.

Apr 2018 – Aug 2022

Full-time study

BSc in Computer Science & Engineering, DUET

Feb 2016 – Aug 2017

Backend Developer (Pre-University)

<redacted>

- **Concurrency & Scale:** Built the reservation engine around a MySQL claim table with row-level locking, preventing double-bookings during flash sales. Took the backend from prototype to production at commercial scale.

02 Projects

github ↗

Custom Raft Consensus Engine · Python 3, network sockets, RPC

- Implemented Raft from the Ongaro–Ousterhout paper (leader election, log replication, heartbeats) and tested safety and partition survival under simulated split-brain faults.

github ↗

Encrypted FUSE Filesystem — sec_fs · C, FUSE3, Linux VFS

- Built a lightweight FUSE3 filesystem in C that transparently encrypts data on write and decrypts on read, designed for secure storage of secrets and keys. Undergraduate thesis at DUET.

03 Skills

Systems & Infra	Linux (systemd, syscalls, VFS, page cache, perf/ftrace, eBPF), Kubernetes, Docker, Nginx, JVM (GC tuning, JFR, JMX, thread dumps)
Distributed & Cloud	Raft, Kafka, Hazelcast, Oracle Coherence, AWS, Oracle Cloud
Networking	TCP/IP, DNS, TLS, socket I/O, Wireshark, F5 BIG-IP LTM
Observability	Datadog (APM, monitors, dashboards), Vector, Falcon LogScale, ELK Stack, Prometheus, Grafana
Languages	Go (goroutines/channels), Python (multiprocessing/C-API), C, C++, Bash, SQL
CI/CD & Automation	Jenkins, Ansible, Git, shell scripting, release pipelines